

Mining Evolving Topics

Web and Social Information Extraction

Project for a.y. 2018/2019

The scope of the project is to identify and trace topics over a temporal interval. A topic can be seen as a set of keywords. Hence, one can see changes in the keyword set of topics throughout the specified time period.

Problem statement: Given $h-l+1$ snapshots, G_1, \dots, G_{h-l+1} , of a temporal research network, identify topics in each snapshot G_i and trace them over the time interval $[l, h]$

Datasets:

DS-1: The keyword co-occurrence

This dataset contains one graph for every year [2000-2018]. Let V be the set of nodes that represent the different **keywords** used by articles in literature: e.g. *machine learning*, *crawling*, *statistical inference*, etc. Let E be the set of edges which represent the relationship of two keywords being used by two different articles. Every edge $e=\{k_i, k_j\}$ is decorated with an ordered dictionary of authors. The dictionary of authors is organized as follows: key-value (a, n) pairs where a represents an author, whereas n is the number of times the author a uses k_i and k_j in his/her articles.

Each row of the dataset is formatted as follows:

$y_q < \text{tab} > k_i < \text{tab} > k_j < \text{tab} > [a_0:n_0, \dots, a_m:n_m] < \text{newline} >$

DS-2: The Co-authorship

This dataset contains one graph for every year [2000-2018]. Let V be the set of nodes representing those authors that have published in the year in question. Let E be the set of edges which depict the relationship of co-authorship: i.e. two nodes a_i and a_j have an edge $e=\{a_i, a_j\}$ if the corresponding authors have published an article together. The weight of an edge e corresponds to the number of collaborations between its two incident nodes.

Each row of the dataset is formatted as follows:

y_q <tab> a_i <tab> a_j <tab> n <newline>

T1: Topic Identification

1. Select *top-k* [$k=5,10,20,100$] (k is the number of generated topics) (according to a certain metric such as pagerank, hits, betweenness, brokerage) keywords in DS-1 **for each year**.
2. For every node in *top-k* apply a Spreading of Influence Algorithm to report the nodes influenced by them in each iteration of the algorithm (similar to Linear Threshold Model or Independent Cascade). *The influenced nodes represent a topic*.
3. Join the produced topics **in a given year** following a merging strategy which takes care of possible overlaps among them.

The students must decide¹:

- which measure should be used to select the top-k starting nodes;
- which is the weight function for the edges;
- which is the threshold function for the nodes;

1. *You can use DS-2 to define your weight and threshold functions.*

T2: Topic Tracing

1. Trace, over the timeline [2000-2018], any topic identified in task **T1**;
2. While analysing the topic temporal/structural behaviour the student must decide if two topics identified in two consecutive years can be merged together;
3. Create a final list of the merged topics;

The students must decide:

- how to determine that a certain topic t_j exposes a similar temporal/structural behaviour of another topic t_i ;